

AI Features - Cost & Usage Report

My Country Mobile | 4 September 2026

Executive summary

All four AI features run on a single model, **Gemini 3.5 Flash-Lite**, using the LLM key each brand already had configured. No new vendor was added and nothing new is provisioned.

Running cost is **\$0.07 to \$1.28 per seat per month** depending on how heavily people use it. AI Writing is effectively free because it only runs when someone clicks; Suggested Replies costs roughly 25 times more because it fires automatically while people type, and accounts for about **96% of the total AI bill**.

At a recommended **\$8-12 per seat per month**, gross margin is approximately **90-98%**.

Read section 5 before acting on this

Three findings materially affect these numbers: the Google account tier is unverified (a free tier would mean zero cost but may allow customer message content to be used for model training), there is no token metering in place today, and there is no per-company spending cap.

1. Model in use

All four features share one model. There is no separate model per feature.

Item	Value
Primary model	gemini-3.5-flash-lite (Google)
Fallback model	gpt-5.4-mini (Azure, eastus2) - used only if Google fails
Configured in	website_settings.ai_details, per brand
New vendor added	No - reuses the brand's existing key
Features covered	AI Writing, Suggested Replies, Conversation Summary, Ask

2. Token usage (measured)

These are not estimates. Each figure was captured from Google's own usageMetadata on real calls against the live service.

Request type	Input	Output	Total
AI Writing - polish, short	78	14	92
AI Writing - elaborate, medium	97	58	155
Suggestions - openers (12-turn chat)	365	72	437
Suggestions - finish sentence (12-turn)	368	126	494
Suggestions - new chat	172	55	227
Suggestions - ceiling, long chat	1,380	216	1,596

Request type	Input	Output	Total
Summary (12-turn)	269	134	403
Summary - ceiling	3,293	39	3,332
Ask (12-turn)	252	8	260
Ask - ceiling	3,276	8	3,284

Why the ceilings matter

Conversation history is capped before it is sent (40 messages for suggestions, 300 for summaries). Cost therefore cannot run away on a long thread: a 320-turn conversation still tops out at roughly 1,600 tokens for suggestions.

3. What it costs

Google's published rate for Gemini 3.5 Flash-Lite is **\$0.30 per million input tokens** and **\$2.50 per million output tokens**.

Cost per 1,000 requests

Feature	Typical	Worst case
AI Writing	\$0.12	\$0.17
Suggested Replies	\$0.36	\$0.95
Conversation Summary	\$0.42	\$1.09
Ask a question	\$0.10	\$1.00

Cost per seat, per month

Usage profile	AI Writing	Suggestions	Both
Light (2 rewrites, 8 suggestions/day)	\$0.005	\$0.06	\$0.07
Medium (3 rewrites, 25/day)	\$0.008	\$0.20	\$0.20
Heavy (8 rewrites, 60/day)	\$0.021	\$0.47	\$0.49
Heavy + long chats (worst)	\$0.021	\$1.26	\$1.28

Total across all 2,068 active seats (23 companies)

Per-request costs above are measured facts. The volumes below are modelled assumptions, since the features are days old and real adoption is not yet known.

Adoption scenario	Monthly	Yearly
10% adopt, light	\$14	\$169
30% adopt, medium	\$127	\$1,521
50% adopt, medium	\$211	\$2,535

Adoption scenario	Monthly	Yearly
100% adopt, medium	\$422	\$5,069
100% adopt, heavy	\$1,018	\$12,217
100% heavy + long chats (worst case)	\$2,647	\$31,758

Where the money actually goes

AI Writing costs about 2 cents per seat per month even for heavy users, because it only runs on a deliberate click. Suggested Replies fires automatically as people type, making it roughly 96% of the total AI spend. If cost ever needs reducing, that is the only lever worth pulling.

4. Recommended customer pricing

Cost is \$0.07-\$1.28 per seat per month; even the worst case stays under \$1.30.

Recommendation: bundle rather than meter. Metering a 20-cent cost creates billing complexity and makes customers ration a feature you want them using. The cost is low and now bounded by the history caps.

Option	Price	Margin	When to use
AI Assist add-on (all four features)	\$8-12 per seat/mo	90-98%	Recommended - clean upsell
Bundle into top tier	+\$5-8 on tier price	85-95%	Drives tier upgrades
Sell features separately	Writing \$3 Suggestions \$6	~95%	Not recommended - better together

Why \$8-12. It sits below the \$15-30 that comparable platforms charge for AI tiers, so it is an easy yes, while still returning roughly 95% margin. Below about \$5 leaves money on the table given the demonstrated value; above about \$15 invites comparison shopping.

Do not price per request. Customers cannot predict their usage, and 36 cents per thousand requests is too small to bill on meaningfully.

5. Three things to confirm before selling this

1. The Google account tier is unverified

Confirming this needs Google Cloud Console access, which was not available. It matters a great deal: on the free tier the cost is zero, but rate limits are hard and **Google may use the data for model training** - a serious problem for customers' private messages. Please confirm the account is on the paid tier before the feature is sold.

2. There is no usage metering at all

Verified: none of the AI services record token counts, and the only AI table written to is the thumbs up/down feedback. Actual spend cannot currently be seen, attributed to a company, or billed per usage. The per-request figures in this report are solid; the volume assumptions are modelled.

3. There is no per-company spending cap

One heavy or abusive tenant can consume without limit. Combined with the absence of metering, this would not surface until the Google invoice arrived.

Suggested next step

Log tokens per request per company. It closes findings 2 and 3 together, and turns the modelled monthly figures in this report into real, invoice-matching ones.

Method. Token counts were captured by intercepting live API calls and reading Google's usageMetadata from the responses, using a representative 12-turn work conversation and a 320-turn conversation to establish the ceilings. Unit prices are Google's published paid-tier rates for Gemini 3.5 Flash-Lite (ai.google.dev/gemini-api/docs/pricing), verified against the official pricing page on 4 September 2026. Seat count (2,068 active users across 23 companies) was read from the production database.